Following the Information Bound from a Search-Window Prior with Sequential Bayesian Experiment Design in Nitrogen-Vacancy Magnetometry

[Authors][1]
[1] *[Affiliation]*
(Dated: November 4, 2026)

[Written last, once the simulation and lab results are in.]

I. INTRODUCTION

Nitrogen-vacancy (N-V) centers in diamond are established sensors of magnetic fields on the nanometer scale [1–3][cite a review, e.g. Degen, Reinhard and Cappellaro, Rev. Mod. Phys. (2017); Rondin et al., Rep. Prog. Phys. (2014)]. In optically detected magnetic resonance (ODMR) magnetometry, the field is a parameter of the spin Hamiltonian of the center and is inferred from a microwave scan: the photoluminescence is recorded at a sequence of microwave drive frequencies within a search window, the range in which the resonance can lie, and the field is read from the position of the resonance. Inferring a Hamiltonian parameter in this way is an instance of Hamiltonian learning [4][cite further Hamiltonian-learning work], in which the experimenter chooses the setting of each measurement, here the drive frequency. Each measurement takes time, so the number of measurements a strategy needs sets both the sensitivity and the bandwidth of the magnetometer. This number matters most where many readings are needed, as in tracking a time-varying field or imaging a field pixel by pixel [cite].

The traditional scan is a sweep, which steps through the window at uniform spacing, and several approaches have been proposed to improve on it. Non-adaptive approaches measure fewer points and reconstruct the spectrum from them, with a low-complexity reconstruction [5] or with compressed sensing, which matches the accuracy of a full scan with a small fraction of its points [6]. Sequential Bayesian experiment design (SBED) is adaptive: it maintains a posterior distribution over the parameters, here represented by a sequential Monte Carlo (SMC) particle cloud, and computes each measurement setting from it by maximizing an expected utility, such as the expected information gain [4]. In ODMR, SBED chooses the next drive frequency and reaches a given precision more than an order of magnitude faster than a sweep [7].

In recent years, machine learning has also been applied to the choice of measurement. Instead of computing each setting from the posterior, these methods learn a policy in advance. Early studies learned policies for adaptive phase estimation [8, 9]. Reinforcement learning has since been used to find the controls that maximize the sensitivity of a quantum sensor [10–12] and to estimate several parameters simultaneously in a photonic experiment [13]. Neural networks trained by an evolution strategy or by reinforcement learning replace the costly optimization of the expected gain in Bayesian design with a fast learned heuristic [14], and model-aware reinforcement learning trains such a policy on a particle-filter posterior, with applications to N-V Ramsey magnetometry [15, 16]. For a wide field range, a Bayesian neural network first narrows the range and a reinforcement-learning agent then refines the estimate [17].

These works present their gains as a multiplicative speedup over a sweep, without relating either strategy to the limit on how quickly any strategy can locate the resonance. Without that comparison it is not known how close the reported strategies come to the best achievable performance, nor how many of their measurements are unavoidable. The speedup itself remains a factor without a mechanism behind it; in particular, its growth as the line occupies a smaller fraction of the window [7] has not been explained.

Here we derive such a limit for the ODMR search and use it to analyze SBED and drive-frequency sweeps for continuous-wave (CW) and pulsed ODMR. The limit follows from Bayesian estimation theory with the search window as the prior [18, 19], and the information it counts is the same quantity that SBED maximizes. It accounts for the whole course of a measurement run: its stages, first finding the resonance, then narrowing in on it and finally refining the estimate, are limits of a single bound rather than separate mechanisms. The same bound gives the speedup over a sweep: for a precise target it is approximately the ratio of the window width to the linewidth, independent of the signal-to-noise ratio, which explains its growth with the window. It also recovers the standard CW-ODMR sensitivity as the refine limit, reached only by a strategy that measures at the steepest slope of the line, and shows that the window limits the bandwidth while the line shape and the signal-to-noise ratio limit the sensitivity. [Simulation and lab findings: how closely SBED follows the bound, and the speedup measured.]

Section II derives the bound from the physical model, together with the speedup, sensitivity and bandwidth that follow from it. Section III presents SBED as an example of a strategy that follows the bound, and describes its particle filter, the sweeps it is compared with and how both are evaluated against the bound. Section IV presents the simulations and the experiment, and Sec. V discusses the limits of the approach.

[Level diagram: $^3A_2$ and ^3E triplets ($m_s = 0, \pm 1$), singlet $|s\rangle$; green pumping, red PL, spin-selective intersystem crossing; microwave transitions at $D \pm \delta_Z$.]

FIG. 1. Five-level model of the N-V center [3, 20]. [Caption to follow the figure.]

[Sequence diagram: CW ODMR (laser and microwaves on together, PL read continuously) and pulsed ODMR (laser pulse to polarize, microwave π-pulse of length $1/(2\Omega)$ with the laser off, laser pulse to read out).]

FIG. 2. Laser and microwave sequences of CW and pulsed ODMR. [Caption to follow the figure.]

II. THEORY

A. Physical model

The negatively charged N-V center is described by five levels [3, 20] (Fig. 1): the spin-triplet ground state $^3A_2$, split into $|0\rangle_g$ and $|\pm 1\rangle_g$ by the spin projection m_s on the N-V axis; the spin-triplet excited state ^3E, with $|0\rangle_e$ and $|\pm 1\rangle_e$; and a metastable spin singlet $|s\rangle$ that stands for the singlet manifold. Green light drives the spin-conserving optical transitions $|m_s\rangle_g \to |m_s\rangle_e$, which decay back by emitting red photoluminescence (PL). From $|\pm 1\rangle_e$, however, the intersystem crossing to the singlet is much faster than from $|0\rangle_e$, and the singlet decays preferentially to $|0\rangle_g$. Two consequences follow. Under continuous illumination the spin is pumped into $m_s = 0$; and a center in $m_s = \pm 1$ spends part of its optical cycles in the dark singlet, so it emits less PL than one in $m_s = 0$. The spin state can therefore be prepared and read out optically.

With a static field B_z along the N-V axis, the ground-state spin ($S = 1$) Hamiltonian is, in units with $h = 1$ (energies in frequency units) [3],

$$H = D\,S_z^2 + \gamma_e B_z S_z + A_\parallel S_z I_z, \tag{1}$$

with $D \approx 2.87$ GHz the zero-field splitting, $\gamma_e \approx 28$ GHz/T the electron gyromagnetic ratio, and $A_\parallel$ the axial hyperfine coupling to the nitrogen nuclear spin I ($|A_\parallel| = 2.15$ MHz for ^{14}N, $I = 1$, and 3.03 MHz for ^{15}N, $I = 1/2$ [21]). The strain splitting $E(S_x^2 - S_y^2)$ is dropped, a small correction when $\gamma_e B_z \gg E$. The eigenstates are the product states $|m_s, m_I\rangle$, and the magnetic-dipole transitions $m_s = 0 \to \pm 1$, which conserve m_I, lie at

$$f_\pm(m_I) = D \pm \left(\delta_Z + A_\parallel m_I\right), \qquad \delta_Z = \gamma_e B_z, \tag{2}$$

where the Zeeman shift δ_Z is the distance of each transition from D: the field moves the two transitions apart, symmetrically about D. The hyperfine term splits each of them further into $2I+1$ lines, one per nuclear projection m_I, spaced by $|A_\parallel|$: a triplet for ^{14}N and a doublet for ^{15}N. Every N-V center contains one of the two isotopes (^{14}N at 99.6% natural abundance, ^{15}N in isotopically enriched or implanted samples), so the hyperfine structure is always present; whether it is resolved depends on the linewidth. We first treat a single line and return to the hyperfine structure below.

An ODMR measurement sets the microwave drive frequency f and records the PL, normalized to its off-resonance level. When f meets a transition, the drive transfers population from $m_s = 0$ to the darker $m_s = \pm 1$ state and the PL dips. We measure the lower transition f_- only (Sec. III D). With the zero-field splitting D treated as a known constant, f_- carries the field, $D - f_- = \delta_Z = \gamma_e B_z$, and an uncertainty in f_- is the same uncertainty in the Zeeman shift.

A drive near f_- couples only $|0\rangle_g$ and $|-1\rangle_g$, and the excited state and the singlet enter only through the rates at which the laser repolarizes the spin and destroys its coherence [22], so the spin is a driven two-level system $\{m_s = 0, m_s = -1\}$. Let the drive have Rabi frequency Ω (Hz, the rate of the population oscillation on resonance) and detuning $\Delta = f - f_-$ from the transition. Off resonance the spin oscillates faster but less deeply, at the generalized Rabi frequency

$$\tilde{\Omega} = \sqrt{\Omega^2 + \Delta^2}.$$

Starting from $m_s = 0$, the probability of finding the spin in $m_s = -1$ after a drive of duration t is

$$P_{-1}(\Delta, t) = \frac{\Omega^2}{\tilde{\Omega}^2}\,\sin^2\!\left(\pi\tilde{\Omega}t\right), \tag{3}$$

and $1 - P_{-1}$ is the probability of finding it still in $m_s = 0$. The prefactor is the depth of the oscillation, the sine its time dependence. Since $m_s = -1$ is the darker state, the PL dip is proportional to P_{-1}. How the drive is applied sets the line shape (Fig. 2).

CW ODMR. In continuous-wave ODMR the laser and the microwave drive act on the spin continuously and at the same time. The laser keeps pumping the spin back into $m_s = 0$ while the drive keeps rotating it away, so every optical reset starts a new Rabi oscillation, and the PL, collected over many such cycles, reports Eq. (3) averaged over the time since the last reset. When the drive acts for many Rabi periods between resets, $\tilde{\Omega}t \gg 1$, the oscillating factor averages to $\frac{1}{2}$. This only renormalizes the amplitude; the line shape is that of the prefactor,

$$\bar{P}_{-1}(\Delta) = \frac{1}{2}\,\frac{\Omega^2}{\Omega^2 + \Delta^2}, \tag{4}$$

a Lorentzian in the detuning Δ of half width Ω. This is the microwave contribution to the width, power broadening: the stronger the drive, the further off resonance it still flips the spin.

Eq. (4) treats the laser only as a reset. Its action on the spin is two rates: each optical cycle pumps the spin back toward $m_s = 0$, at a rate Γ_p, and destroys the phase between $m_s = 0$ and $m_s = -1$, at a rate Γ_c. These are

a longitudinal and a transverse relaxation rate, $1/T_1$ and $1/T_2$ of the two-level system. The steady state of the Bloch equations with these rates and a drive of angular Rabi frequency $\Omega_R = 2\pi\Omega$ is [22]

$$P_{-1}^{\rm ss}(\Delta) = \frac{1}{2}\,\frac{\Omega_R^2\,\Gamma_c/\Gamma_p}{(2\pi\Delta)^2 + \Gamma_c^2 + \Omega_R^2\,\Gamma_c/\Gamma_p}, \tag{5}$$

which reduces to Eq. (4) when the two rates are equal and small against Ω_R. It is still a Lorentzian, now of half width

$$\gamma = \frac{1}{2\pi}\sqrt{\Gamma_c^2 + \Omega_R^2\,\frac{\Gamma_c}{\Gamma_p}}. \tag{6}$$

At weak drive the width is set by the laser, through Γ_c; at strong drive by the microwave power. Both rates are set by how often the center completes an optical cycle, and that rate saturates. In terms of the optical saturation parameter $s = I_{\rm opt}/I_{\rm sat}$, the laser intensity in units of the intensity that saturates the optical transition, the cycling rate grows linearly with s at low power and levels off at high power, where it is limited by the lifetime of the excited state rather than by the laser: the center cannot be excited again before it has decayed. Both rates follow it, $\Gamma_{p,c} = \Gamma_{p,c}^{\infty}\,s/(1+s)$ [22], with $\Gamma_{p,c}^{\infty}$ their values at full saturation. The depth of the line, Eq. (5) at $\Delta = 0$, is

$$P_{-1}^{\rm ss}(0) = \frac{1}{2}\,\frac{\Omega_R^2}{\Omega_R^2 + \Gamma_p\Gamma_c}:$$

it grows with the microwave power, toward $\frac{1}{2}$, and falls with the laser power, which repolarizes the spin before the drive can flip it. Raising the laser power therefore both widens and flattens the line. Whatever its origin, the line is one Lorentzian dip: the normalized PL expected at drive frequency f is

$$\mu(f) = 1 - C\,\frac{\gamma^2}{\gamma^2 + (f - f_-)^2}, \tag{7}$$

with contrast C, the depth of the dip at $f = f_-$, half width γ of Eq. (6) and full width at half maximum $\Gamma = 2\gamma$. Optical and microwave broadening are not modeled separately: they enter only through the single linewidth Γ, which is inferred with the line.

Pulsed ODMR removes the laser from the width of the line [22]. In CW ODMR the laser that reads the spin out also broadens the line through Γ_c, and narrowing it means lowering both powers and with them the contrast and the PL. Pulsed ODMR separates the two in time: a laser pulse polarizes the spin into $m_s = 0$, the laser is switched off while a single microwave π-pulse of duration $t_\pi = 1/(2\Omega)$ is applied, and a second laser pulse reads the spin out (Fig. 2). Eq. (3) is then evaluated at the one time t_π instead of being averaged. The pulse length is chosen so that on resonance, where the spin oscillates at Ω, it is fully inverted: $\pi\Omega t_\pi = \pi/2$. Off resonance it oscillates at $\tilde\Omega$ instead. Measuring $\tilde\Omega$ in units of Ω,

$$r \equiv \frac{\tilde\Omega}{\Omega} = \sqrt{1 + \Delta^2/\Omega^2}\ \geq 1,$$

the argument of the sine in Eq. (3) at $t = t_\pi$ is $\pi\tilde\Omega t_\pi = \pi r/2$ and the prefactor is $\Omega^2/\tilde\Omega^2 = 1/r^2$, so

$$P_{-1}^{\pi}(\Delta) = \frac{\sin^2(\pi r/2)}{r^2}. \tag{8}$$

On resonance $r = 1$ and the inversion is complete, $P_{-1}^{\pi} = 1$. Off resonance the line is not a Lorentzian: for $|\Delta| \gg \Omega$, $r \approx |\Delta|/\Omega$ and $P_{-1}^{\pi} \approx \sin^2(\pi\Delta t_\pi)/(2\Delta t_\pi)^2$, the power spectrum of a rectangular pulse of length t_π: far from resonance the spin responds to the frequency content of the pulse. The half width of Eq. (8) is $0.80\,\Omega = 0.40/t_\pi$, so the line narrows as the pulse gets longer. This is the Fourier (time–bandwidth) limit: a pulse of duration t_π contains frequencies spread over about $1/t_\pi$, so it cannot address the spin more selectively than that. The line vanishes wherever $\sin(\pi r/2) = 0$, i.e. $r = 2, 4, \ldots$, at $|\Delta| = \sqrt{3}\,\Omega, \sqrt{15}\,\Omega, \ldots$: there the spin completes a whole number of full Rabi cycles and returns to $m_s = 0$. Between the zeros lie side lobes that decay as $1/r^2$; the first, near $r = 3$ at $|\Delta| \approx 2.7\,\Omega$, is 12% of the depth. Side lobes are a hazard for a search: a lobe is a dip where there is no resonance.

The inhomogeneous dephasing time T_2^* sets how far this narrowing can go. The transition frequency of the spin is not sharp: slowly fluctuating local fields, mostly from the bath of ^{13}C nuclear spins, shift it by a random offset ξ that is fixed during one pulse but changes from one repetition of the sequence to the next. A measurement averages many repetitions, each driven at the detuning $\Delta - \xi$ instead of Δ, so the measured line is Eq. (8) averaged over the distribution of ξ, i.e. convolved with it. For a Gaussian distribution of standard deviation σ_ξ,

$$\bar P_{-1}^{\pi}(\Delta) = \int P_{-1}^{\pi}(\Delta - \xi)\,\frac{e^{-\xi^2/2\sigma_\xi^2}}{\sqrt{2\pi}\,\sigma_\xi}\,d\xi. \tag{9}$$

The width σ_ξ is what T_2^* measures. A free spin precessing at the offset ξ picks up the phase $2\pi\xi t$, and the average of $e^{2\pi i\xi t}$ over the Gaussian is the coherence envelope $e^{-2\pi^2\sigma_\xi^2 t^2}$. T_2^* is defined by writing this envelope as $e^{-(t/T_2^*)^2}$, so $\sigma_\xi = 1/(\sqrt{2}\,\pi T_2^*)$. For $t_\pi \ll T_2^*$ the spread is small against Ω and the line is that of Eq. (8). For $t_\pi \gtrsim T_2^*$ the Gaussian dominates: the zeros and side lobes are washed out, the line stops narrowing at a width $\sim 1/T_2^*$, and the depth falls because the pulse no longer inverts every frequency class [22]. The best pulsed line is therefore reached at $t_\pi \sim T_2^*$, a width set by the spin itself rather than by the laser or microwave power.

Both pulses take a finite time, and the spin evolves during each of them. While the microwave pulse is on, the laser is off and the spin relaxes only through T_1, which is much longer than t_π, and dephases through the spread already included in Eq. (9). The readout is different: the laser that reads the spin out also repolarizes it, so the PL difference between $m_s = 0$ and $m_s = -1$ lasts only for the repolarization time τ_p. Taking the repolarization as a single exponential, a spin flipped with probability

P_{-1} emits, a time t into the readout pulse, PL at the rate $\Phi_0\,[1 - c\,P_{-1}\,e^{-t/\tau_p}]$, where Φ_0 is the rate of $m_s = 0$ and c the fractional PL drop of $m_s = -1$ at the start of the readout. The recorded PL is this rate integrated over a readout window of length t_r, normalized to the same integral with no flip,

$$\begin{aligned}\mu &= 1 - \frac{1}{t_r}\int_0^{t_r} c\,P_{-1}\,e^{-t/\tau_p}\,dt \\ &= 1 - c\,\frac{\tau_p}{t_r}\big(1 - e^{-t_r/\tau_p}\big)\,\bar{P}^{\pi}_{-1}(\Delta),\end{aligned} \tag{10}$$

with $P_{-1} = \bar{P}^{\pi}_{-1}$ of Eq. (9). The integration leaves the line shape unchanged and sets only its depth: the contrast is $C = c\,(\tau_p/t_r)(1 - e^{-t_r/\tau_p})$. A longer window collects more photons, $N_{\mathrm{ph}} \propto t_r$, but dilutes the contrast, so the single-shot signal-to-noise ratio $C\sqrt{N_{\mathrm{ph}}} \propto (1 - e^{-t_r/\tau_p})/\sqrt{t_r/\tau_p}$ is largest at $t_r \approx 1.26\,\tau_p$.

Hyperfine structure. Each nuclear projection m_I gives its own copy of the line, shifted by $A_\parallel m_I$ (Eq. (2)) and weighted by the population of that m_I, so the observed spectrum is a sum of $2I + 1$ lines. Two cases reduce it to the single line of Eq. (7). Near the excited-state level anticrossing the nitrogen nuclear spin is polarized into a single m_I [23]: only one line remains, and its fixed offset $A_\parallel m_I$ is absorbed into f_-. When the line is much wider than the hyperfine spacing, $\Gamma \gg |A_\parallel|$, the lines merge into one. [Whether the main text keeps the single line or treats the resolved hyperfine structure. Where the hyperfine-resolved runs are used, and how f_- is defined for them (center line or m_I-averaged).]

The second transition. The spectrum is mirror symmetric about D, so the drive window is placed below D and holds the lower transition only. The tail of the upper one, at $f_+ = f_- + 2\delta_Z$, adds $C\gamma^2/[\gamma^2 + (f - f_+)^2]$ inside the window: at the line center a fraction $\gamma^2/(\gamma^2 + 4\delta_Z^2)$ of the contrast C, i.e. 20% at $\delta_Z = \gamma$ and below 1% for $\delta_Z \gtrsim 5\gamma$. Equation (7) assumes $\delta_Z \gg \gamma$. Small Zeeman shifts, where the tails of the two lines overlap, are a separate problem and outside the scope of this work.

f_- is the parameter of interest. The width Γ and the contrast C are nuisance parameters: they are inferred jointly with f_- but only matter through their effect on it. The noise level σ (below) is inferred as well.

Noise. Each reading of the normalized PL is noisy. Its floor is photon shot noise: a reading that collects N_{ph} photons on average off resonance has a relative standard deviation $N_{\mathrm{ph}}^{-1/2}$, and slightly less inside the dip, where fewer photons arrive. Technical noise (laser intensity, microwave power, collection drift) adds to it. We model a single reading, a shot, as

$$y = \mu(f) + \epsilon, \qquad \epsilon \sim \mathcal{N}(0, \sigma^2), \tag{11}$$

with one $\sigma \geq N_{\mathrm{ph}}^{-1/2}$ at every drive frequency; this neglects the small reduction of the shot noise inside the dip, of relative order of the contrast C. A measurement is the mean of k shots at the same f, so its noise is $\sigma/\sqrt{k}$. The signal-to-noise ratio of a single shot depends on where it is taken: it is the drop in PL at f over the noise,

$$\mathrm{SNR}(f) = \frac{1 - \mu(f)}{\sigma} = \frac{C}{\sigma}\,\frac{\gamma^2}{\gamma^2 + (f - f_-)^2}. \tag{12}$$

It is largest on resonance, where the drop is the contrast C,

$$\mathrm{SNR} \equiv \mathrm{SNR}(f_-) = C/\sigma, \tag{13}$$

the contrast over the noise, the single number by which a line is labeled. All SNR values in this paper follow Eq. (13). It neglects the mirror line: its tail raises the on-resonance depth to $C[1 + \gamma^2/(\gamma^2 + 4\delta_Z^2)]$, so C/σ is the $\delta_Z \gg \gamma$ limit of the on-resonance SNR.

The prior knowledge of the experimentalist is an interval of width W that holds the resonance, $f_- \in [f_{\min}, f_{\min} + W]$, and the drive frequency is scanned over at least this interval. The ratio W/Γ, the inverse of the fraction of the window the line occupies, sets how hard the search is. [State the value of W used in the simulations.]

The model above predicts what a measurement records for given parameters. The rest of this section turns the question around: given the measurements, how well can f_- be known, how much can each measurement add, and where should the next one be taken. These questions are answered first without reference to any particular estimator, through the Fisher information and the prior, and then in the Bayesian terms used by the strategies.

B. Information of a measurement

Choosing where to measure is a problem of experiment design. We use its standard notation: the design $x \equiv f$ is the drive frequency at which a measurement is taken, the outcome y is the normalized PL it records, the parameter of interest $\theta \equiv f_-$ is the transition frequency, and $\phi = (\Gamma, C)$ are the nuisance parameters. The expected normalized PL $\mu(x; \theta, \phi)$ of Eq. (7), now written as a function of the design and the parameters, is the mean outcome, and Eq. (11) gives the spread about it.

How much one measurement at x tells about θ is its Fisher information. For a Gaussian outcome of variance σ^2/k it depends only on how fast the predicted outcome moves when θ moves, the slope g:

$$I(x) = \frac{k\,g(x)^2}{\sigma^2}, \qquad g(x) = \frac{\partial \mu(x; \theta, \phi)}{\partial \theta}. \tag{14}$$

A shift $\delta\theta$ moves the mean outcome by $g\,\delta\theta$, and it is resolved once that exceeds the noise $\sigma/\sqrt{k}$; $I(x)$ is the inverse square of the smallest shift resolved. It grows linearly with k, and at the shot-noise limit, $\sigma^2 = 1/N_{\mathrm{ph}}$, with the total number of photons kN_{ph} collected at that drive frequency, which no averaging scheme can beat.

For the Lorentzian of Eq. (7), written in terms of the detuning $u = x - \theta$ of the measurement from the resonance,

$$g(x) = \frac{\partial \mu}{\partial \theta} = -\frac{\partial \mu}{\partial u} = -\frac{2C\gamma^2 u}{(\gamma^2 + u^2)^2},$$

which vanishes on resonance, where the dip is flat, and on the baseline far from it: a measurement at the bottom of the dip or far outside it says almost nothing about where the dip is. $|g|$ is largest on the flanks, at $|u| = \gamma/\sqrt{3}$, where the dip is 3/4 deep, with

$$g^* \equiv \max_x |g(x)| = \frac{9C}{4\sqrt{3}\,\Gamma} \approx 1.30\,\frac{C}{\Gamma}.$$

The best possible information per measurement is therefore

$$I^* = \max_x I(x) = \frac{k\,g^{*2}}{\sigma^2} = \frac{27}{16}\,k\,\frac{\mathrm{SNR}^2}{\Gamma^2}, \qquad (15)$$

with $\mathrm{SNR} = C/\sigma$ of Eq. (13). I^* assumes the nuisance parameters (width, contrast) are known; with free nuisances it is an upper bound on the usable information. The nuisances are inferred along with θ but are not the subject of this work.

Measurements are independent given θ, so their Fisher information adds: n measurements at $x_1, \ldots, x_n$ carry

$$J_n = \sum_{i=1}^{n} I(x_i) \;\le\; n\,I^*,$$

with equality only when every measurement is taken at the steepest slope of the true line. By the Cramér–Rao bound [24], an unbiased estimate of θ from them has a standard deviation of at least $J_n^{-1/2} \ge (nI^*)^{-1/2}$.

C. Prior information: the search window

Fisher information measures knowledge in inverse variance, Hz^{-2}: a Gaussian estimate of standard deviation σ_θ is worth $J = \sigma_\theta^{-2}$, and combining independent Gaussian estimates adds their inverse variances. The same currency measures what is known *before* the first measurement. If the experimentalist knows only that the resonance lies somewhere in the window of width W, the state of knowledge is a uniform prior $p(\theta) = 1/W$ on the window. Its spread is its standard deviation σ_0; with v the distance from the center of the window,

$$\sigma_0^2 = \int_{-W/2}^{W/2} \frac{v^2}{W}\,dv = \frac{W^2}{12}, \qquad J_0 \equiv \sigma_0^{-2} = \frac{12}{W^2}. \quad (16)$$

J_0 is the information the window holds: it is what a Gaussian estimate with the same variance would be worth. A narrower window starts with more, information that no measurement had to pay for. A uniform prior itself has no finite Fisher information, since its density jumps at the edges of the window; taking the inverse variance J_0 in its place is a convention, exact for a Gaussian prior.

D. Van Trees bound

Because the prior carries information, the relevant bound is not the Cramér–Rao bound, which holds for unbiased estimators and ignores J_0, but the van Trees inequality [18, 19]. It holds for any estimator, biased or not, with the error averaged over the prior, and it adds the information of the prior to that of the measurements:

$$\sigma_\theta(n) \;\ge\; \Big(J_0 + \sum_{i=1}^{n} I(x_i)\Big)^{-1/2} \;\ge\; \sigma_{\mathrm{vT}}(n) = \Big(\frac{12}{W^2} + nI^*\Big)^{-1/2}, \qquad (17)$$

where $\sigma_\theta(n)$ is the root-mean-square error after n measurements and each $I(x_i)$ is the information of measurement i, both averaged over the prior. The first inequality is the van Trees inequality with J_0 in place of the prior's Fisher information, and the second uses $I(x) \le I^*$. Since the Fisher information of any prior is at least its inverse variance, this substitution can only raise the curve: for the uniform window σ_{vT} is a reference rather than a strict lower bound, and for a Gaussian prior it is exact.

The bound constrains the mean squared error averaged over the prior, whatever the shape of the posterior. For the posterior mean that error is the expected posterior variance, so the posterior standard deviation can be compared with σ_{vT} from the first measurement on, as an average over runs. A Gaussian posterior is needed only to read σ_θ as a 68% interval and for the exact form of the information gain below.

The bound has two limits, set by how the information of the measurements compares with that of the prior. While $nI^* \ll J_0$, it stays at the prior, $\sigma_{\mathrm{vT}} \approx \sigma_0 = W/\sqrt{12}$: the measurements have not yet added to what the window told. An estimate in this limit leans on the prior and is biased toward its center, which is why its error may lie below the unbiased Cramér–Rao bound $(nI^*)^{-1/2}$. Once $nI^* \gg J_0$, the prior is washed out and $\sigma_{\mathrm{vT}} \approx (nI^*)^{-1/2}$, the unbiased Cramér–Rao bound, which falls with slope $-\frac{1}{2}$ on a log–log scale. That the posterior then forgets the prior is not shown here but taken from the Bernstein–von Mises theorem [25]: for a parameter of any finite dimension, here (θ, ϕ), and under regularity conditions, the posterior approaches, as the data accumulate, a Gaussian centered near the true value whose covariance is the inverse of the Fisher information matrix $\mathbf{J}_n$ of all the parameters. Its mean is then asymptotically unbiased and attains the Cramér–Rao bound. The marginal posterior variance of θ is the $\theta\theta$ element of $\mathbf{J}_n^{-1}$, which is at least $1/J_n$, the value with the nuisances known; the two are equal when θ is orthogonal to the nuisances, i.e. when the off-diagonal elements of $\mathbf{J}_n$ vanish. For the Lorentzian this holds when the measurements are placed symmetrically on the two flanks, since $\partial\mu/\partial\theta$ is odd in u while $\partial\mu/\partial\Gamma$ and $\partial\mu/\partial C$ are even. The theorem is proved for measurements fixed in advance; here each design depends on the previous outcomes, so we use it as a guide to the data limit rather than as a result. The two limits cross at $n_c = J_0/I^*$.

We name the regimes after what a strategy does in them: *search* in the prior limit, where it has yet to find the dip; *refine* in the data limit, where it measures on the flanks of a dip it has located; and *focus* in the crossover between them, where it narrows down from the window to the dip.

If every measurement carried I^*, the crossover would come after $n_c = J_0/I^* = \frac{64}{9}\,(\Gamma/W)^2/(k\,\mathrm{SNR}^2)$ measurements, far below one for any realistic window ($n_c \approx 2\times10^{-4}$ for $W = 20\,\Gamma$ at SNR = 10). It comes much later because I^* is reached only once the dip has been found: a measurement on the baseline carries $I(x) \approx 0$, so the sum in Eq. (17) stays near J_0 while the strategy searches. What such a measurement contributes instead is *elimination*: it shrinks the prior itself.

a. Elimination model. Consider a strategy that sweeps the window at spacing Δx (e.g. $\Delta x = \Gamma/3$) until it finds the dip and then measures at the steepest slope. Let the SNR be so high that a single measurement taken within the line, $|x-\theta| \lesssim \gamma$, shows a drop that cannot be mistaken for noise, so the first measurement that lands in the dip detects it. After n measurements without a dip, the resonance is known only to lie in the part of the window not yet excluded, and the posterior is uniform there, so

$$\sigma(n) = \frac{W - n\Delta x}{\sqrt{12}}, \qquad n < n_d, \tag{18}$$

where $n_d \le W/\Delta x$ is the step that finds the dip (on average $W/(2\Delta x)$ over a uniformly placed resonance). In information terms $J(n) = 12/(W-n\Delta x)^2$: the gain per step is not constant but grows as the remaining width shrinks. Detection localizes the resonance to about one grid step, a uniform interval of width Δx worth $12/\Delta x^2$, after which each measurement adds I^*:

$$\sigma(n) = \Big(\frac{12}{\Delta x^2} + (n-n_d)\,I^*\Big)^{-1/2}, \qquad n \ge n_d. \tag{19}$$

Eqs. (18)–(19) reproduce both limits of Eq. (17): $\sigma = \sigma_0$ at $n = 0$ and $\sigma \to (nI^*)^{-1/2}$ for $n \gg n_d$. For a wide window the three regimes become visible (Fig. 3). The search is a plateau: Eq. (18) falls linearly, which on a log–log axis looks flat until about half the window is excluded, for a number of steps that grows with $W/\Delta x$. The focus is a cliff steeper than $-\frac{1}{2}$, where detection collapses σ from about the window to about Δx in a single step. It ends once σ is small compared with the line, so that the strategy knows where the steepest slope is. The refine regime is the $-\frac{1}{2}$ tail. A narrow window ($W \lesssim 3.5\,\Gamma$, $\sigma_0 \lesssim \Gamma$) has no plateau. Search, focus and refine are therefore not separate mechanisms but the regimes of one bound.

The elimination model only illustrates the shape of the profile. At low SNR one measurement no longer reveals the dip and n_d grows, but Eq. (17) holds at any SNR and needs no detection model; the runs are read against it alone.

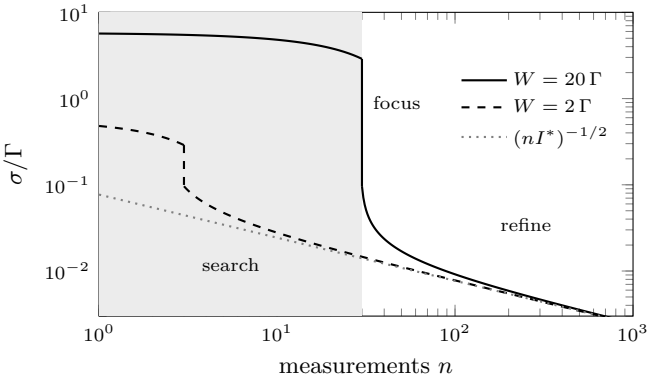


FIG. 3. Schematic of the three regimes: the elimination model, Eqs. (18)–(19), at illustrative values (SNR 10, $k = 1$, $\Delta x = \Gamma/3$, dip found halfway through the grid). Solid: wide window, $W = 20\,\Gamma$; dashed: narrow window, $W = 2\,\Gamma$; dotted: unbiased Cramér–Rao bound $(nI^*)^{-1/2}$. The wide window shows a search plateau (shaded), a focus cliff at detection (vertical segment) and the $-\frac{1}{2}$ refine tail; the narrow window shows no plateau. Both start near the prior limit σ_0 of the van Trees bound and tend to its data limit, the unbiased Cramér–Rao bound, at large n. Simulated and measured runs are shown against the bound in Sec. IV.

E. Information gained per measurement

The bound counts Fisher information, the most a measurement can contribute. What a measurement actually contributes is how far it moves the posterior: the Kullback–Leibler divergence from the posterior before it to the posterior after it. Before the outcome is known, the expectation of this divergence is the expected information gain [26],

$$\begin{aligned}\mathrm{EIG}(x) &= \mathbb{E}_{y|x}\Big[D_{\mathrm{KL}}\big(p(\theta,\phi \mid y,x)\,\|\,p(\theta,\phi)\big)\Big] = I(\theta,\phi;y \mid x)\\ &= H[y \mid x] - \mathbb{E}_{\theta,\phi}H[y \mid \theta,\phi,x],\end{aligned} \tag{20}$$

the mutual information between the parameters and the outcome, with $p(\theta,\phi)$ the current posterior. Summed over a run, the gains are what the measurements added to the knowledge, and greedy Bayesian experiment design takes at every step the x that maximizes Eq. (20).

The second term is fixed by the noise model: the noise is Gaussian with the same variance σ^2/k at every drive frequency (Sec. II A), so $H[y \mid \theta,\phi,x] = \frac{1}{2}\ln(2\pi e\,\sigma^2/k)$ for every x. All the dependence on x is in the first term, the entropy of the predictive distribution of y, whose spread is the noise plus the spread $\sigma^2_{\mathrm{pred}}(x)$ of the predicted signal $\mu(x;\theta,\phi)$ over the posterior. For a given variance the Gaussian has the largest entropy [27], so $H[y \mid x] \le \frac{1}{2}\ln 2\pi e\,(\sigma^2_{\mathrm{pred}}(x) + \sigma^2/k)$ and

$$\mathrm{EIG}(x) \le \tfrac{1}{2}\ln\Big(1 + \frac{k\,\sigma^2_{\mathrm{pred}}(x)}{\sigma^2}\Big), \tag{21}$$

with equality when the predictive distribution of $\mu(x;\theta,\phi)$ is Gaussian. Entropies here are in nats (natural

logarithms). Whether it is depends on the posterior of θ, not on the noise: the noise is Gaussian by construction, the posterior only becomes Gaussian with enough data.

a. When the variance is exact: refine. In the data limit of Sec. II D the posterior of θ is close to a Gaussian of variance σ_θ^2 (equal to $1/J_n$ when θ is orthogonal to the nuisances). Once it is also narrow compared with the line, μ is linear in θ over it, $\mu(x;\theta) \approx \mu(x;\hat\theta)+g(x)(\theta-\hat\theta)$, so the predicted signal is Gaussian too, with $\sigma^2_{\text{pred}}(x) = g(x)^2\sigma_\theta^2$, and Eq. (21) holds with equality:

$$\text{EIG}(x) = \tfrac{1}{2}\ln\left(1 + I(x)\,\sigma_\theta^2\right) = \tfrac{1}{2}\ln\frac{J_n + I(x)}{J_n}. \qquad (22)$$

Maximizing the EIG is then maximizing $I(x)$, i.e. measuring at the steepest slope, and the gains telescope: $\sum_n \text{EIG}_n = \frac{1}{2}\ln(J_n/J_0) = \ln(\sigma_0/\sigma_n)$, the logarithm of the factor by which the uncertainty has narrowed. The information the acquisition rule maximizes and the information budget of Sec. II D are the same quantity. A strategy that measured every point at the steepest slope of the true line would reach the data limit of Eq. (17); greedy EIG approaches it once the posterior is narrow compared with the line.

Eq. (22) thus needs two conditions: the posterior has a single mode, with no appreciable mass left elsewhere in the window, and that mode is narrow compared with the line, $\sigma_\theta \ll \gamma$. Appendix B shows that the deviations from a Gaussian, and the error of the linearization, are both controlled by the ratio σ_θ/γ, and that leftover mass far from the dip dominates the posterior variance unless it is smaller than about $(\sigma_\theta/W)^2$. Both conditions are met in refine and not at the start of focus, so Eq. (22) describes the refine regime, while in focus Eq. (21) holds only as a bound. [Results: when the SBED posterior becomes Gaussian, and its calibration in refine.]

b. When it is not: search and focus. Before the dip is found the posterior of θ is spread over the window, and the predicted signal at x is two-valued: with probability p, the posterior mass that would put the dip at x, it is about $1-C$, otherwise 1. Then $\sigma^2_{\text{pred}} = p(1-p)C^2$, while the true mutual information is at most the binary entropy $H_b(p) \le \ln 2$. The variance form keeps growing with SNR where the true gain saturates, so it overvalues search candidates relative to refine candidates; both peak at $p = \frac{1}{2}$, so they rank search candidates alike. Here greedy EIG is only myopically optimal, which is where the strategies of Sec. III B differ.

The same two-valued picture holds in focus, with a different ranking. Once a peak has formed at the dip, a small residual mass p far from it gives a predictive variance $p(1-p)C^2$ at its location, large at high SNR even for small p, while the information it holds is at most $H_b(p)$, small for small p. The variance form therefore keeps spending measurements to remove such distant mass. This reduces the posterior variance, which the tail dominates and which is the quantity Eq. (17) constrains, rather than the posterior entropy: in focus the variance rule acts as a design for the mean squared error, while a rule built on the information gain would refine the peak sooner and accept the tail. Whether this explains why focus costs SBED far more measurements than the bound is a hypothesis we have not tested. [Test: share of focus measurements placed more than a few FWHM from the peak, compared with the residual mass there.]

Bayesian SBED (Sec. III B) evaluates Eq. (21) at every candidate drive frequency on the particles of the particle filter of Sec. III A, which learns σ and uses its posterior mean. [Mention computing the exact mixture entropy from the particles, or leave out?]

F. Speedup over a drive-frequency sweep

The same budget predicts the speedup over a conventional sweep. An adaptive strategy in the refine regime measures at the steepest slope and gains, per measurement, Eq. (15),

$$I^* = \frac{27}{16}\,\frac{k\,C^2}{\sigma^2\,\Gamma^2}.$$

A sweep spreads its measurements uniformly over the window, so its information per measurement is the window average of Eq. (14),

$$\bar I = \frac{1}{W}\int I(x)\,dx = \frac{k}{\sigma^2 W}\int g(x)^2\,dx = \frac{\pi}{2}\,\frac{k\,C^2}{\sigma^2\,\Gamma\,W}, \qquad (23)$$

using $\int g^2\,du = \pi C^2/(4\gamma) = \pi C^2/(2\Gamma)$ for the Lorentzian, valid for $W \gg \Gamma$. Both are $\propto kC^2/\sigma^2$; they differ in the width factor, $1/\Gamma^2$ against $1/(\Gamma W)$, because the sweep spends only a fraction of order Γ/W of its measurements on the line.

The sweep therefore accumulates information at the rate $\bar I$ from its first measurement, $\sigma \approx (J_0 + n\bar I)^{-1/2}$, with no plateau and no cliff, as observed by Dushenko *et al.* [7]. An adaptive strategy pays the search cost n_d first and then gains I^* per measurement. Let σ_{tgt} (Hz) be the target precision, the standard deviation of f_- at which a run is considered done. The number of measurements to reach it is

$$\begin{aligned} N_{\text{sweep}}(\sigma_{\text{tgt}}) &= \frac{\sigma_{\text{tgt}}^{-2} - J_0}{\bar I}, \\ N_{\text{ad}}(\sigma_{\text{tgt}}) &= n_d + \frac{\sigma_{\text{tgt}}^{-2} - 12/\Delta x^2}{I^*}, \end{aligned} \qquad (24)$$

where the adaptive count follows the elimination model: n_d measurements to detect the dip, after which the information starts at $12/\Delta x^2$ and grows by I^* per measurement until it reaches σ_{tgt}^{-2}. The speedup is $R(\sigma_{\text{tgt}}) = N_{\text{sweep}}/N_{\text{ad}}$. For a precise target the refine terms dominate and

$$\begin{aligned} R_\infty = \frac{I^*}{\bar I} &= \frac{27}{8\pi}\,\frac{W}{\Gamma} \\ &\approx 1.07\,\frac{W}{\Gamma} \approx \frac{W}{\Gamma}\,: \end{aligned} \qquad (25)$$

the speedup is the inverse occupancy of the window, the number of linewidths that fit in it, which Dushenko *et al.* [7] identified qualitatively; Eq. (25) gives its prefactor. It is independent of the SNR, because both strategies gain information $\propto \mathrm{SNR}^2$. For a loose target the search cost n_d dominates N_{ad} and $R(\sigma_{\mathrm{tgt}}) < R_\infty$, so the full dependence on window, linewidth, SNR (through n_d at low SNR) and target is given by Eq. (24). A sweep cannot approach the van Trees bound at any n: its rate $\bar{I}$ is fixed by the window, while the bound's data limit requires every measurement at the steepest slope. [Check Eq. (25) against the speedup reported in Ref. [7] for their window and line structure (a resolved triplet, not one Lorentzian).]

G. Sensitivity and bandwidth

Sensitivity and bandwidth follow from the same budget once measurements are converted to time. Let one measurement take $t_m = k\,t_s$ (k shots of duration t_s), so that n measurements take $T = nt_m$. With D known, $D - f_- = \delta_Z = \gamma_e B_z$, so the field uncertainty is $\sigma_B = \sigma_{f_-}/\gamma_e$.

a. Sensitivity (refine). In the refine regime every measurement adds I^*, so after a time T

$$\sigma_{f_-}(T) = (nI^*)^{-1/2} = \Big(\frac{T}{t_m}\,I^*\Big)^{-1/2} \quad\Longrightarrow\quad \sigma_{f_-}\sqrt{T} = \sqrt{\frac{t_m}{I^*}}.$$

The uncertainty falls as $\sigma_{f_-} \propto T^{-1/2}$, so the product $\sigma_{f_-}\sqrt{T}$ does not depend on T and characterizes the strategy rather than the length of the run. Converted to field, this constant is the sensitivity,

$$\eta = \sigma_B\sqrt{T} = \frac{1}{\gamma_e}\sqrt{\frac{t_m}{I^*}} = \frac{1}{\gamma_e}\sqrt{\frac{16\,k\,t_s\,\Gamma^2}{27\,k\,\mathrm{SNR}^2}} = \frac{4}{3\sqrt{3}}\,\frac{\Gamma\sqrt{t_s}}{\gamma_e\,\mathrm{SNR}}, \tag{26}$$

with $\mathrm{SNR} = C/\sigma$ the single-shot SNR of Eq. (13). The number of shots k cancels as well: averaging k shots into one measurement costs k times the time and gives k times the information. The prefactor $4/(3\sqrt{3})$ is that of the standard CW-ODMR sensitivity of a Lorentzian line [22]: the familiar formula is the refine limit of the budget, reached only by a strategy that measures at the steepest slope. A strategy whose refine measurements are not all at the steepest slope gains less per measurement: if they carry on average a fraction $q \le 1$ of I^*, I^* is replaced by qI^* in Eq. (26), and its sensitivity is worse by $1/\sqrt{q}$, i.e. $\eta/\sqrt{q}$.

b. Bandwidth (whole run). A field reading is useful once its error reaches the target σ_{tgt}; the time needed, $T(\sigma_{\mathrm{tgt}})$, sets the bandwidth $1/T(\sigma_{\mathrm{tgt}})$. It is t_m times the number of measurements $N_{\mathrm{ad}}(\sigma_{\mathrm{tgt}})$ of Eq. (24), n_d to find the dip and the rest to refine it to σ_{tgt}:

$$\begin{aligned} T(\sigma_{\mathrm{tgt}}) &= t_m\Big[n_d + \frac{\sigma_{\mathrm{tgt}}^{-2} - 12/\Delta x^2}{I^*}\Big], \\ \eta_{\mathrm{eff}}(\sigma_{\mathrm{tgt}}) &= \frac{\sigma_{\mathrm{tgt}}\sqrt{T(\sigma_{\mathrm{tgt}})}}{\gamma_e}, \end{aligned} \tag{27}$$

where η_{eff} is the sensitivity reached at that target, η with σ_{tgt} in place of σ_{f_-}. For $\sigma_{\mathrm{tgt}} \to 0$ the refine term dominates, $T \approx t_m\sigma_{\mathrm{tgt}}^{-2}/I^*$, and $\eta_{\mathrm{eff}} \to \eta$ of Eq. (26). For a loose target the search term dominates, $T \approx n_d t_m$: the bandwidth cannot exceed about $1/(n_d t_m) \sim 2\Delta x/(Wt_m)$, whatever the SNR, and η_{eff} is worse than η. The window, i.e. the prior, therefore limits the bandwidth, while the line shape and SNR limit the sensitivity. Only the total integration time $T = nk\,t_s$ enters the bound, not how it is divided into measurements: n measurements of k shots carry the same information as nk measurements of one shot at the same drive frequency. Neglecting the overhead of switching the drive frequency, the cost of a run is therefore its total number of shots nk, the integration time in units of t_s, and strategies that group their shots differently are compared on that axis.

All closed-form results of this section are for the Lorentzian line of Eq. (7). For pulsed ODMR the same quantities follow, numerically, with Eq. (10) in place of Eq. (7). [Which ODMR mode the results focus on, once the simulations and the experiment are complete; where pulsed goes (a Theory subsection for the information of the pulsed line, and its own Results).]

III. METHODS

A. Bayesian inference

Every strategy runs the same loop: it chooses a drive frequency x_n, takes one measurement y_n there, updates the posterior over the line parameters with it, and stops when the stop rule of Sec. III B fires. The strategies differ only in the first step; the posterior and its update, described here, are shared. The unknowns are those of Sec. II B, the parameter of interest θ and the nuisances ϕ, together with the noise level σ, which is treated separately below.

The posterior over (θ, ϕ) is a weighted cloud of particles $\{(\theta_i, \phi_i), w_i\}_{i=1}^{N_p}$, updated by sequential Monte Carlo [28]. Each particle is one hypothesis for the line. It holds the Zeeman shift $\delta_{Z,i}$, so $\theta_i = D - \delta_{Z,i}$ with D fixed (the simulations measure the mirror transition $D + \delta_{Z,i}$, which is equivalent by symmetry), and the nuisances: the linewidth itself, as the half width $\gamma = \Gamma/2$ of Eq. (7), and the contrast C. The mechanisms that set the width, Eq. (6), are not inferred; only the width they produce is. For the pseudo-Voigt line the width is held as its Lorentzian half width and the standard deviation of its Gaussian part.

The noise is Gaussian, Eq. (11), as it is for the photon numbers of an ODMR shot, but its level σ is not known in advance: technical noise adds to the shot noise by an amount that depends on the setup and can drift. σ is therefore inferred with the line, but not as a particle dimension. Given a particle's line, the Gaussian likelihood is conjugate to an inverse-gamma prior on σ^2, so each particle carries that posterior in closed form, updated from its own residuals, and σ^2 is integrated out exactly: a Rao–Blackwellized particle filter with per-particle sufficient statistics [29, 30]. A measurement then reweights each particle by a Student-t predictive, the Gaussian likelihood averaged over that particle's noise posterior. Integrating out matters: a likelihood with a plug-in estimate of σ is largest for a particle whose σ happens to equal one residual, so resampling would favor such particles and drive σ to its lower bound, while the Student-t tails do not reward that. Old residuals are discounted, as in the variance discounting of dynamic linear models [31], so the noise posterior can follow a slow drift of σ.

When the weight has concentrated on a few particles, the cloud is resampled systematically [32] and jittered with the Liu–West kernel [33], which shrinks every particle towards the mean of the cloud and adds Gaussian noise, so that the mean and covariance of the cloud are preserved. Two safeguards depart from it: a variance floor that decays with the number of measurements, so early rejuvenations can still widen the cloud, and reflection of particles that leave the prior bounds. Each particle's noise posterior is then recomputed from the full history at its new position. The initial cloud is drawn from the prior [align with Sec. II C]. The formulas and settings are in Appendix A.

The particle filter reports the posterior mean and the weighted particle standard deviation,

$$\hat{\theta}_N = \sum_i w_i \theta_i, \qquad \hat{\sigma}_N = \Big[\sum_i w_i (\theta_i - \hat{\theta}_N)^2 \Big]^{1/2}, \quad (28)$$

the claimed uncertainty whose meaning Sec. III C discusses.

B. Measurement strategies

We compare two families of strategies. The sweep family holds the strategies an experimentalist would use without adaptive design: plain sweeps at several grid spacings, and Smart-Sweep, a sweep refined inside the dip once it is seen. The question is whether the best member of that family can match Bayesian SBED. All strategies infer with the particle filter of Sec. III A, so the comparison isolates the choice of measurements from the inference, and $\hat{\sigma}_N$ means the same for all of them. All run under the same stop rule and the same measurement budget: a run stops once the claimed uncertainty of every parameter has stayed below its ceiling for several consecutive measurements and the dip test below finds a dip.

a. Dip test. The test asks where the measurements, not the particles, show a dip. Each measurement is scored against a baseline taken from the upper part of the data, in units of the noise from the noise posterior. Contiguous stretches of measured drive frequencies, about as wide as a plausible line, get a boxcar score, and a stretch is significant when its score passes a family-wise false-alarm threshold over all stretches tested. Each connected run of significant stretches is one dip.

b. Sweep family. SimpleSweep measures a uniform grid over the window, a few points per FWHM of the prior-mean line, and repeats whole passes; the spacing is the one setting a sweep has, and several are run. Smart-Sweep sweeps the window coarsely until the dip test finds a dip, then finely inside the focus (below), and then measures only at the steepest-slope points of the posterior-mean line, once the posterior expects them to carry nearly the optimal information. Its coarse search uses the same grid points, and the same measured values, as the reference sweep of that repeat.

c. Bayesian SBED. SBED is the particle filter choosing each x by Eq. (21), maximized over candidate drive frequencies concentrated on the flanks of the posterior-mean line and on the dips found. The candidates are restricted to the focus: one section per detected dip, padded by one half width on each side, which widens at once when a new dip appears and narrows only after a delay. A strategy that has piled its measurements onto a suspected dip cannot be contradicted elsewhere, so with a small probability, largest while much of the window is unmeasured, SBED instead measures the middle of the widest unmeasured gap, a falsification probe.

C. Simulation setup and evaluation

a. What the claimed uncertainty is. The particle filter reports $\hat{\sigma}_N$, Eq. (28): a moment of the posterior, defined at every N. The van Trees bound, Eq. (17), limits the mean squared error averaged over the prior, and for the posterior mean that error equals the expected posterior variance, $\mathbb{E}[(\hat{\theta}_N - \theta)^2] = \mathbb{E}[\hat{\sigma}_N^2]$, when the posterior is correct. Comparing $\hat{\sigma}_N$ with the bound is therefore valid from the first measurement, as an average over runs and provided the posterior is calibrated. What fails before the posterior is Gaussian (Sec. II E) is the usual reading of $\hat{\sigma}_N$: with a uniform window or several live candidates it is not a 68% interval, and the posterior entropy is below $\frac{1}{2}\ln(2\pi e\,\hat{\sigma}_N^2)$, so information gains, per-step gains as variance ratios, and coverage at Gaussian quantiles are only meaningful in the data regime. We therefore (i) compare $\hat{\sigma}_N$ with the bound as a run average at all N, (ii) test coverage, the true error against $\hat{\sigma}_N$, only once the posterior is Gaussian, or with quantiles of the particle distribution itself, and (iii) keep this separate from miscalibration of the particle filter, which makes $\hat{\sigma}_N$ too small at any N.

b. Use cases. A use case is one simulated signal at one noise level, set by the SNR, the drive window, the Gaussian share of the width and the isotope (Appendix A). Every metric is computed within one use case, over repeats whose true signal is the same for every strategy, so strategies are compared run by run; several use cases are summarized by the median over them.

c. Metrics. The accuracy at the end of a run is not a quality metric: every run stops once its claimed uncertainty is below the ceiling, so its final error mostly reflects the ceiling. We report instead whether the claimed uncertainty is honest, through its coverage, the share of runs whose true error lies within $\hat{\sigma}_N$ and within $2\hat{\sigma}_N$; how close each strategy comes to the bound of Eq. (17), through the ratio of the bound to the mean squared true error at each N; and the cost, the number of measurements after which the true error stays below a target. Each run is also read against the bound through its information rate, the growth of $\hat{\sigma}_N^{-2}$ per measurement in units of I^*, Eq. (15); a rate above 1 means the claimed uncertainty fell faster than the data allow. Measurements to stop are compared only between strategies whose claimed uncertainty is equally honest.

D. Experimental setup

We measure the lower transition f_- only: in our setup it needs less microwave power than f_+, which keeps the heating of the stage low and the amplifier within its linear range. [Sample, optics, microwave delivery, detection and the per-shot integration time t_s.]

IV. RESULTS

A. Simulations

[(a) Claimed σ and true error against Eqs. (17)–(19) for a narrow and a wide window. (b) Information rate and prior share. (c) Cost vs SNR and window. (d) Calibration of the claimed σ.]

B. Experiment

[Placeholder until the lab data are in.]

V. DISCUSSION

[Limits: the claimed σ is overconfident; the simulated prior is miscalibrated; Gaussian noise only; I^* assumes known nuisances; low-SNR detection model.]

VI. CONCLUSION

[Conclusion.]

Appendix A: Algorithm settings

a. Particle filter. The cloud has $N_p = 1000$ particles. Every dimension is scaled to $[0, 1]$ by its prior bounds, so one resampling kernel serves all of them; reported values are mapped back to Hz. The cloud is resampled when the effective sample size $(\sum_i w_i^2)^{-1}$ falls below $0.2\,N_p$, systematically, and then rejuvenated with the Liu–West kernel,

$$\psi_i' \sim \mathcal{N}\big(a\psi_i + (1-a)\bar{\psi},\ (1-a^2)\Sigma\big), \qquad a = 0.98, \quad \text{(A1)}$$

with $\psi_i = (\theta_i, \phi_i)$ the whole particle and $\bar{\psi}$, Σ the weighted mean and covariance before resampling. The diagonal of Σ is floored at $(0.01\,e^{-n/25})^2$ of each unit range after n measurements, and a particle that leaves the unit cube is reflected back into it.

b. Noise model. Each particle carries an inverse-gamma posterior on σ^2 with shape α, the same for every particle because its update does not depend on the residuals, and scale β_i. The prior is $\alpha_0 = 10$, $\beta_0 = \alpha_0\sigma_{\mathrm{nom}}^2$, with σ_{nom} the geometric mean of the prior range of σ. Before each measurement both are discounted by $d = 0.99$; a measurement y_n, the mean of k shots at x_n, then adds $\frac{1}{2}$ to α and $\frac{1}{2}k\,r_{n,i}^2$ to β_i, with residual $r_{n,i} = y_n - \mu(x_n; \theta_i, \phi_i)$, so that after n measurements $\beta_i = d^n\beta_0 + \frac{1}{2}\sum_{t<n} d^{n-1-t} k\, r_{t,i}^2$. Old residuals are thus forgotten with a memory of about $1/(1-d) = 100$ measurements, and α saturates at $0.5/(1-d) = 50$, about 100 degrees of freedom. With σ^2 integrated out, the measurement reweights the particles by the Student-t predictive

$$w_i \ \propto\ w_i\, t_{2\alpha}\Big(y_n;\ \mu(x_n;\theta_i,\phi_i),\ \frac{\beta_i}{k\alpha}\Big), \qquad \text{(A2)}$$

taken with α, β_i before the update. After a rejuvenation each β_i is recomputed from the full history at the particle's new position.

c. Initial cloud. [The prior on the Zeeman shift δ_Z is being changed to uniform over the range the truth is drawn from (now a truncated Gaussian, offset from the truth, plus 10% of the particles uniform over the bounds); state it here once Sec. II C is fixed.]

d. Stop rule and budget. The ceilings on the claimed uncertainty are 100 kHz for θ, 500 kHz for each width and 0.01 for the contrast; they must hold for 8 consecutive measurements, together with a dip found by the dip test, which is re-run at most every 10 measurements. Every strategy has the same budget of 2500 measurements; a sweep's total is capped in whole passes. A run that reaches the budget without stopping is censored and counted at the budget, as a lower bound on its cost.

e. Dip test. Each measurement gets the score $z_n = (b - y_n)/s$. The baseline b is the 70th percentile of the y_n minus its expected offset under pure noise, $t_\nu^{-1}(0.7)\,s$, and s and ν are the scale and degrees of freedom of the noise posterior. Every contiguous stretch of measured drive frequencies whose span lies in the prior range of the linewidth gets the boxcar score $\sum z_n/\sqrt{n}$, over its n measurements. A stretch is significant when its score exceeds $t_\nu^{-1}(1 - p_{\mathrm{FA}}/N_{\mathrm{tests}})$, the Bonferroni threshold for a family-wise false-alarm probability $p_{\mathrm{FA}} = 0.1$ over the N_{tests} stretches tested.

f. Sweep family. Γ_0 is the FWHM of the prior-mean line. SimpleSweep measures a uniform grid of spacing Γ_0/p, $p = 2, 3, 4$. Smart-Sweep sweeps at $\Gamma_0/3$ until the dip test finds a dip, then at $\Gamma_0/10$ inside the focus, and then measures only at the maxima of $g(x)^2$ at the posterior mean, cycling between them, once the posterior expects each of those points to carry at least 90% of the information of the true optimum.

g. Bayesian SBED. After every resampling SBED places 800 candidate drive frequencies at the quantiles of a density with Gaussian kernels on the flanks, ±HWHM, of the posterior-mean line, kernels on the dips found and a flat share, thinned to a spacing of at least 100 kHz. Eq. (21) is evaluated on a 500-particle subsample, with σ^2 the posterior mean of the noise. The candidates are taken in blocks of 64, and the block maxima compete through a softmax with a temperature of 0.01 nats, so the choice is near-greedy. Each focus section is a dip's extent padded by one posterior-mean HWHM on each side. A new dip widens the focus at once; narrowing waits 8 measurements and must shrink the focus by at least 5%. The focus is re-decided after every resampling. The falsification probe is made with probability $0.02 + 0.08\,g$, where g is the widest unmeasured gap as a fraction of the window.

h. Use cases. The simulated signals are pseudo-Voigt dips of total FWHM 2 MHz, on a grid of the SNR, the window (1.6, 2, 3 and 5 times the width one dip group covers), the Gaussian share of the width (0, 0.15, 0.3) and the isotope (unresolved hyperfine lines, ^{14}N or ^{15}N). The code labels the SNR levels 1, 2, 5, 10, 20 against a contrast of 0.25 [convert to C/σ of Eq. (13): the code's contrast is $2C$]. [Number of repeats per use case.]

Appendix B: Approach of the posterior to a Gaussian

Eq. (22) and the reading of σ_θ as a 68% interval assume a Gaussian posterior. This appendix states when that holds for the line of Eq. (7). There are two separate questions: whether the posterior has a single peak, and whether that peak is Gaussian. We take the nuisances as known; they add dimensions but do not change the scaling.

a. Shape of a single peak. Let $L(\theta) = \ln p(\theta) + \sum_i \ln p(y_i \mid \theta, x_i)$ be the log posterior and $\hat\theta$ its maximum. Inside the window the uniform prior is flat and only the likelihood contributes. Expanding about the maximum (Laplace's method [25]),

$$L(\theta) = L(\hat\theta) - \frac{J}{2}(\theta-\hat\theta)^2 + \frac{L'''}{6}(\theta-\hat\theta)^3 + \frac{L''''}{24}(\theta-\hat\theta)^4 + \dots, \tag{B1}$$

with $J = -L''(\hat\theta)$, whose expectation is the Fisher information J_n of Eq. (17). Keeping only the quadratic term gives a Gaussian of variance $\sigma_\theta^2 = 1/J$. Over the posterior, $|\theta - \hat\theta| \sim \sigma_\theta$, so the higher terms change the log density by relative amounts of order $L'''\sigma_\theta^3$ and $L''''\sigma_\theta^4$; the first sets the asymmetry of the peak, the second the weight of its tails.

For Gaussian noise, $\ln p(y \mid \theta, x) = -k\,[y - \mu(x;\theta)]^2/2\sigma^2$ up to a constant. Averaged over the noise, the derivatives are

$$\mathbb{E}[L''] = -\sum_i \frac{k\,\mu_i'^2}{\sigma^2} = -J_n, \qquad \mathbb{E}[L'''] = -3\sum_i \frac{k\,\mu_i'\mu_i''}{\sigma^2},$$

where primes are derivatives with respect to θ at x_i, so $\mu_i' = g(x_i)$. The line changes on the scale of its half width γ, so each further derivative brings a factor of order $1/\gamma$: $\mu'' \sim \mu'/\gamma$, hence $|L'''| \lesssim J_n/\gamma$ and $|L''''| \lesssim J_n/\gamma^2$. The corrections in Eq. (B1) are then of relative order

$$L'''\sigma_\theta^3 \sim \frac{\sigma_\theta}{\gamma}, \qquad L''''\sigma_\theta^4 \sim \frac{\sigma_\theta^2}{\gamma^2}. \tag{B2}$$

A single peak is Gaussian once it is narrow compared with the line, $\sigma_\theta \ll \gamma$. This is an upper estimate: at the steepest-slope points $\mu'' = 0$ (there $|g|$ is at its maximum), so for measurements placed there the cubic term vanishes on average and the corrections start at order σ_θ^2/γ^2.

In the data limit, $J_n \approx nI^*$ with I^* of Eq. (15), and with $\Gamma = 2\gamma$,

$$\frac{\sigma_\theta}{\gamma} \approx \frac{8}{3\sqrt{3}}\,\frac{1}{\mathrm{SNR}\sqrt{nk}} \approx \frac{1.5}{\mathrm{SNR}\sqrt{nk}},$$

so at an SNR of order 10 a few shots on the flanks already make the peak Gaussian.

b. Linearization of the signal. The same ratio controls the step from a Gaussian posterior to a Gaussian prediction used in Eq. (22). Expanding the signal over the posterior, $\mu(x;\theta) = \mu(x;\hat\theta) + g(x)(\theta-\hat\theta) + \frac{1}{2}\mu''(x)(\theta-\hat\theta)^2 + \dots$, the quadratic term relative to the linear one is $\mu''\sigma_\theta/2g \sim \sigma_\theta/\gamma$. Eq. (22) therefore holds up to relative corrections of order σ_θ/γ, and of order σ_θ^2/γ^2 at the steepest slope.

c. Leftover mass. The expansion describes one peak. Early in focus the posterior is instead a mixture: a peak of weight $1 - p$ and width σ_θ at the dip, and a remainder of weight p spread over the part of the window not yet excluded, of width up to W. The variance of the mixture is

$$\mathrm{Var}(\theta) \approx (1-p)\,\sigma_\theta^2 + p\,s_{\mathrm{r}}^2,$$

where $s_r^2 \lesssim W^2/12$ is the mean squared distance of the remainder from the peak. The remainder dominates once $p \gtrsim \sigma_\theta^2/s_r^2$, i.e. roughly $p \gtrsim 12(\sigma_\theta/W)^2$, a very small weight: for $W = 20\,\Gamma$ and $\sigma_\theta = 0.1\,\Gamma$ it is 3×10^{-4}. Until the remainder has fallen below this, the posterior standard deviation reports the remainder, not the error of the peak, and the posterior is far from Gaussian: a narrow peak with wide tails. Its entropy, by contrast, is changed by the remainder only by about $H_b(p)+p\ln(s_r/\sigma_\theta)$, which is small for small p. This is the difference between the variance and the information gain discussed in Sec. II E.

The remainder is removed only by measuring where it would place the dip. A measurement within about γ of such a location, when the dip is not there, lowers the log weight of that part of the remainder on average by the Kullback–Leibler divergence between the two predicted outcomes, $kC^2/2\sigma^2 = k\,\mathrm{SNR}^2/2$. Lowering a weight p to below $12(\sigma_\theta/W)^2$ takes, at each location, about $2\ln[pW^2/12\sigma_\theta^2]/(k\,\mathrm{SNR}^2)$ measurements, and at least one; a remainder spread over the whole window has of order W/Γ such locations. At high SNR the posterior therefore becomes Gaussian after about one pass over the remaining window, a cost that grows with W/Γ like that of the search. [Compare with when the SBED posterior becomes Gaussian in the runs (Results).]

[1] G. Balasubramanian, I. Y. Chan, R. Kolesov, M. Al-Hmoud, J. Tisler, C. Shin, C. Kim, A. Wojcik, P. R. Hemmer, A. Krueger, T. Hanke, A. Leitenstorfer, R. Bratschitsch, F. Jelezko, and J. Wrachtrup, Nature **455**, 648 (2008).

[2] J. R. Maze, P. L. Stanwix, J. S. Hodges, S. Hong, J. M. Taylor, P. Cappellaro, L. Jiang, M. V. G. Dutt, E. Togan, A. S. Zibrov, A. Yacoby, R. L. Walsworth, and M. D. Lukin, Nature **455**, 644 (2008).

[3] M. W. Doherty, N. B. Manson, P. Delaney, F. Jelezko, J. Wrachtrup, and L. C. L. Hollenberg, Physics Reports **528**, 1 (2013).

[4] C. E. Granade, C. Ferrie, N. Wiebe, and D. G. Cory, New Journal of Physics **14**, 103013 (2012).

[5] S. Luo, Z. Zhu, L. Zhang, P. Chen, and R. Lin, Optical and Quantum Electronics **55**, 1004 (2023).

[6] G. Haim, C. Mullarkey, J. Howell, and N. Bar-Gill, Compressed sensing enabled high-bandwidth and large dynamic range magnetic sensing (2025), arXiv:2502.06070 [quant-ph].

[7] S. Dushenko, K. Ambal, and R. D. McMichael, Physical Review Applied **14**, 054036 (2020).

[8] A. Hentschel and B. C. Sanders, Physical Review Letters **104**, 063603 (2010).

[9] A. Lumino, E. Polino, A. S. Rab, G. Milani, N. Spagnolo, N. Wiebe, and F. Sciarrino, Physical Review Applied **10**, 044033 (2018).

[10] H. Xu, J. Li, L. Liu, Y. Wang, H. Yuan, and X. Wang, npj Quantum Information **5**, 82 (2019).

[11] J. Schuff, L. J. Fiderer, and D. Braun, New Journal of Physics **22**, 035001 (2020).

[12] T. Xiao, J. Fan, and G. Zeng, npj Quantum Information **8**, 2 (2022).

[13] V. Cimini, M. Valeri, E. Polino, S. Piacentini, F. Ceccarelli, G. Corrielli, N. Spagnolo, R. Osellame, and F. Sciarrino, Advanced Photonics **5**, 016005 (2023).

[14] L. J. Fiderer, J. Schuff, and D. Braun, PRX Quantum **2**, 020303 (2021).

[15] F. Belliardo, F. Zoratti, F. Marquardt, and V. Giovannetti, Quantum **8**, 1555 (2024).

[16] F. Belliardo, F. Zoratti, and V. Giovannetti, Physical Review A **109**, 062609 (2024).

[17] S. Guo, J. Liu, T. Le, and H. Dai, Scientific Reports 10.1038/s41598-026-76354-7 (2026).

[18] H. L. Van Trees, *Detection, Estimation, and Modulation Theory, Part I* (Wiley, New York, 1968).

[19] R. D. Gill and B. Y. Levit, Bernoulli **1**, 59 (1995).

[20] N. B. Manson, J. P. Harrison, and M. J. Sellars, Physical Review B **74**, 104303 (2006).

[21] S. Felton, A. M. Edmonds, M. E. Newton, P. M. Martineau, D. Fisher, D. J. Twitchen, and J. M. Baker, Physical Review B **79**, 075203 (2009).

[22] A. Dréau, M. Lesik, L. Rondin, P. Spinicelli, O. Arcizet, J.-F. Roch, and V. Jacques, Physical Review B **84**, 195204 (2011).

[23] V. Jacques, P. Neumann, J. Beck, M. Markham, D. Twitchen, J. Meijer, F. Kaiser, G. Balasubramanian, F. Jelezko, and J. Wrachtrup, Physical Review Letters **102**, 057403 (2009).

[24] H. Cramér, Princeton University Press (1946).

[25] A. W. van der Vaart, *Asymptotic Statistics* (Cambridge University Press, Cambridge, 1998).

[26] D. V. Lindley, The Annals of Mathematical Statistics **27**, 986 (1956).

[27] T. M. Cover and J. A. Thomas, *Elements of Information Theory*, 2nd ed. (Wiley, Hoboken, NJ, 2006).

[28] A. Doucet, N. de Freitas, and N. Gordon, eds., *Sequential Monte Carlo Methods in Practice* (Springer, New York, NY, 2001).

[29] G. Storvik, IEEE Transactions on Signal Processing **50**, 281 (2002).

[30] C. M. Carvalho, M. S. Johannes, H. F. Lopes, and N. G. Polson, Statistical Science **25**, 88 (2010).

[31] M. West and J. Harrison, *Bayesian Forecasting and Dynamic Models*, 2nd ed. (Springer, New York, 1997).

[32] G. Kitagawa, Journal of Computational and Graphical Statistics **5**, 1 (1996).

[33] J. Liu and M. West, in *Sequential Monte Carlo Methods in Practice*, edited by A. Doucet, N. de Freitas, and N. Gordon (Springer, New York, NY, 2001) pp. 197–223.